

SUMMARY

Software Engineer with 6 years shipping production data platforms and ELT pipelines across high-growth B2B SaaS companies, with deep hands-on experience in Snowflake: Snowpipe, Snowpark, UDFs, Stored Procedures, Integrations, compute optimization, warehouse-native architecture. Strong data engineering foundation — pipeline design, data modeling, and cloud data warehousing — paired with software engineering fundamentals: Python, containerized deployments, IaC, and CI/CD

TECHNICAL SKILLS

Languages: Python, SQL, Java, Unix Shell Scripting

Data Platforms & Orchestration: Snowflake, Airflow, dbt, Spark, Databricks, PostgreSQL, Fivetran

Cloud & Infra: Docker, AWS, GCP, Terraform, Docker, Kubernetes, CI/CD, Git, TDD, Refactoring

Monitoring & Observability: Grafana, Prometheus (exposure), PagerDuty, Monte Carlo, CloudWatch, Datadog

Data Architecture: ETL/ELT, CDC, Data Modeling (Star/Snowflake Schema), OLAP, Data Governance, Knowledge Graphs

EDUCATION

MSc Advanced Computing | King's College London, UK | Expected Graduation: September 2026

Thesis: Data Protection Index - Multi-agent LLM pipeline for global privacy law assessment and enforcement scoring

BSc Industrial Engineering | Universidad de Concepción, Chile | 2015–2021 | *Maximum Distinction*

Thesis: Optimization algorithms for large-scale resource allocation (Google Machine Reassignment Problem)

WORK EXPERIENCE

Software Engineer (Data), DataChef, Remote (10/2025-Present)

- Architected a serverless data platform on AWS (Lambda, SQS, Aurora) with infrastructure as code and automated CI/CD, generating \$100K ARR for an enterprise customer. (*Python, Django, AWS, Terraform, Docker*)

Data Engineer II, Front, Chile (03/2025–10/2025)

- Built an event-driven ingestion engine using AWS Lambda and Snowpipe for real-time app event processing with change-data-capture patterns and optimised incremental loading, reducing Snowflake compute costs by 95% (*Python, SQL, AWS, Terraform, Datadog*)
- Designed an automated data quality framework using dbt and Airflow with anomaly detection, profiling checks, and staging environment validation across production pipelines, surfacing 30% storage waste. (*Python, dbt, Airflow*)

Data Engineer II, PagerDuty, Chile (04/2023-03/2025)

- Achieved \$1.5M reduction in annual ad spend with Digital Ads Performance Platform, a distributed data system processing 2TB+ daily from 10+ sources with Airflow orchestration, Monte Carlo quality monitoring, and PagerDuty observability dashboards (*Python, SQL, Snowflake, Spark, Airflow*)
- Led migration of an 80TB data warehouse to FedRAMP-compliant AWS infrastructure, implementing automated ETL workflows, data governance policies, access controls, and end-to-end validation with zero downtime. (*AWS, Terraform, Airflow, Snowflake, Databricks*)
- Built a self-service analytics platform with custom ingestion connectors for 15+ sources (Salesforce, Jira, product databases) via REST APIs, designing star schema models that unified customer lifecycle, pipeline, and revenue metrics into trusted BI dashboards used by 30+ business stakeholders. (*Python, Airflow, Snowflake, AWS*)

Lead Backend Engineer, Datawheel, Chile (10/2020–03/2023)

- Drove \$30K revenue with OEC Premium, a data platform on GCP with OLAP architecture and REST/SQL APIs serving 1,000+ external users, maintaining comprehensive API documentation enabling self-service data access for economic researchers globally (*Python, SQL, FastAPI, ClickHouse, GCP, BigQuery*)
 - Enabled \$25K revenue increase with DataMexico Platform, architecting OLAP data infrastructure with established coding standards and comprehensive documentation for ETL workflows, serving custom aggregations to government partners through scalable APIs (*Python, GCP, OLAP, ETL*)
-

PROJECTS

- Architecting a medallion data platform on AWS (S3, Lambda, Glue, RDS, Neptune) ingesting and reconciling 119+ data sources — including sanctions lists, BIS Entity Lists, GLEIF corporate registries, and OpenSanctions feeds — into a knowledge graph of corporate ownership hierarchies and sanctions exposure, enabling automated risk scoring for clients operating under export control regulations. (*Python, AWS Neptune, RDF/SPARQL, dbt, Airflow*)